

Athreya Anand

US Citizen

San Francisco, CA | +1 (661) 414-2869 | athreyaanand.com | athreyanand@gmail.com | /in/athreyaanand

Skills: Python, Java TypeScript, SQL · PyTorch, LangChain, RAG, MCP, agents, evals, fine-tuning · Anthropic, OpenAI, Gemini (VertexAI) · Kubernetes (GKE), AWS, Spanner, pgvector

Experience

Google – New York, NY

October 2022 – Present

Software Engineer, Tech Lead

Partner Intelligence Engine (PIEngine) – part of Google's internal A.I. Incubator

- Led the design and build from scratch of an agentic workflow orchestration platform — configurable nodes (forms, human review, AI summarization, fine-tuned classifiers, publishing) on Borg + Spanner.
- Powers 20+ partner workflows across Google TV, Home, and Civics; 10,000+ reviews monthly; projected to save millions annually in extended-workforce cost.
- Integrated Gemini summarization and fine-tuned classifiers that pre-score each review; human-reviewed labels feed back to progressively expand auto-decisioning.

AMigo – Internal Hackathon Winner

- Won internal hackathon with AMigo — a RAG assistant for Account Managers tracking 100s of products; replaced relationship-driven routing with data-backed recommendations.
- Built nightly Vertex AI embedding pipeline over each product's TDDs / PRDs / one-pagers; Gemini-powered chat recommends the right team/product for inbound requests.

Google Ventures

- Built ML pipeline for deal-success prediction over 1.7M angel investments (\$613B aggregate); benchmarked GNNs, deep nets, regressions, and tree ensembles.
- Shipped random forest with engineered features at ~80% accuracy; deployed to GV partners to inform deal selection.

Mission Control

- Led migration of legacy CMS to Google Kubernetes Engine (GKE) and database to CloudSQL.
- Refactored codebase and architecture — page-load seconds → milliseconds; deploy cycle days → ~2 hours.

Cedar – New York, NY

May 2021 – September 2022

Software Engineer

- Engineered healthcare data ingestion pipelines (ETLs) processing \$1B+ in charges for Providence Health (Cedar's largest client, 9th-largest US health system) and several other SBO Hospitals.
- Compressed client onboarding from months to weeks via centralized integration tooling; built data validation notebooks adopted across the org.

Amazon (AWS) – Cupertino, CA

May 2019 - August 2019

Machine Learning SWE Intern — AWS Inferentia

- Built per-chip metrics and observability for AWS Inferentia (Amazon's custom ML inference accelerator); surfaced utilization and performance telemetry through the EC2 stack.
- Fleet dashboards adopted by the Inferentia team to tune the chip's runtime, identify bottlenecks, and validate inference workload migrations.

Projects

Viki – tryviki.com

- Production SaaS that turns any codebase into a shared atlas — visual canvas for humans, MCP server for AI agents. Semantic graph built from AST (Tree-sitter, 15+ languages) + OpenAI embeddings + pgvector + graph ML models.
- Measured impact on coding agents: 20.8× fewer LLM tokens per task (196,958 → 9,478 median) and +23.4 pp on SWE-bench Lite (33.3% → 56.7%, p = 0.020) at equal answer quality across paired runs and 20 cents per ingestion.

Habitability – habitability.app

- Cross-platform social habit-tracking app for iOS and Android — solo-built in Flutter; built around social accountability
- Grew to 4,000 MAU within ~5 months of launch.

Education

Georgia Institute of Technology – Atlanta, GA

M.S. Computer Science/Artificial Intelligence & Machine Learning; GPA: 4.0

May 2021

B.S. Computer Science; Honors Program; Cum. GPA: 3.75; CS GPA 4.0

May 2020